



Writing resilient microservices with Dapr

Shubham Sharma, Microsoft
@shubham1172

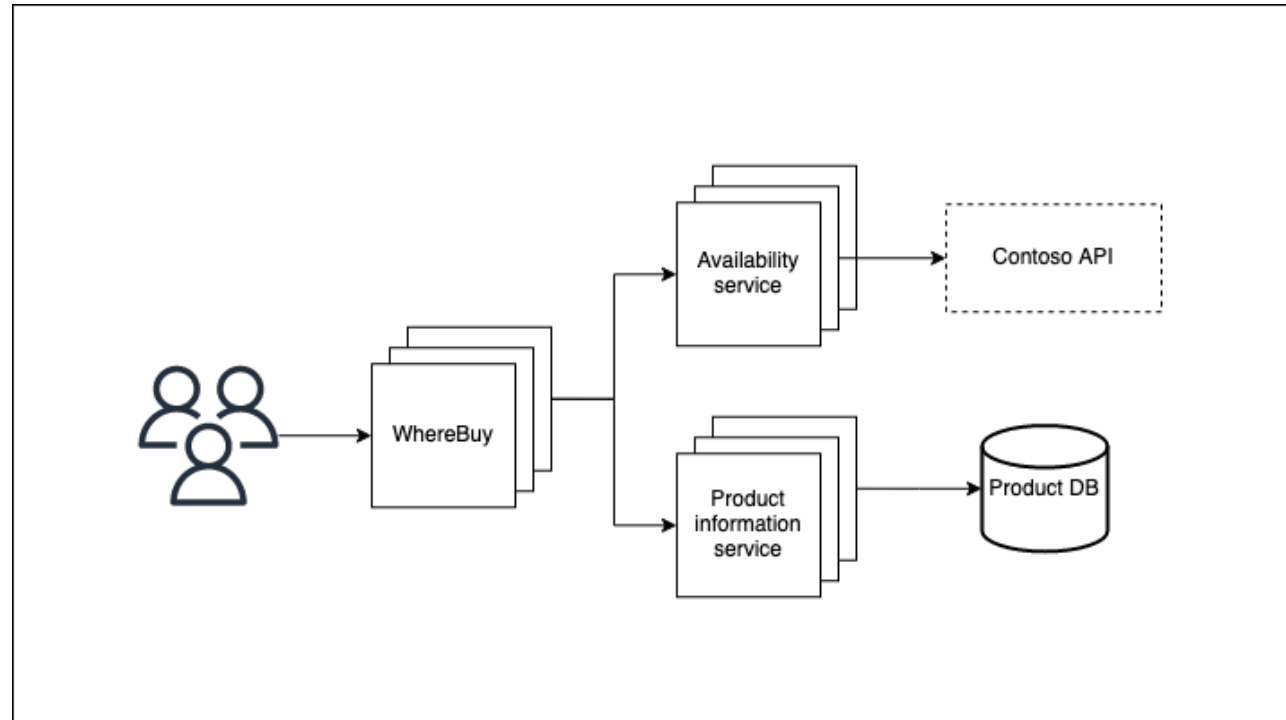
What's inside

- Chaos in the real world
- Resiliency with Dapr
- Dapr-izing your applications

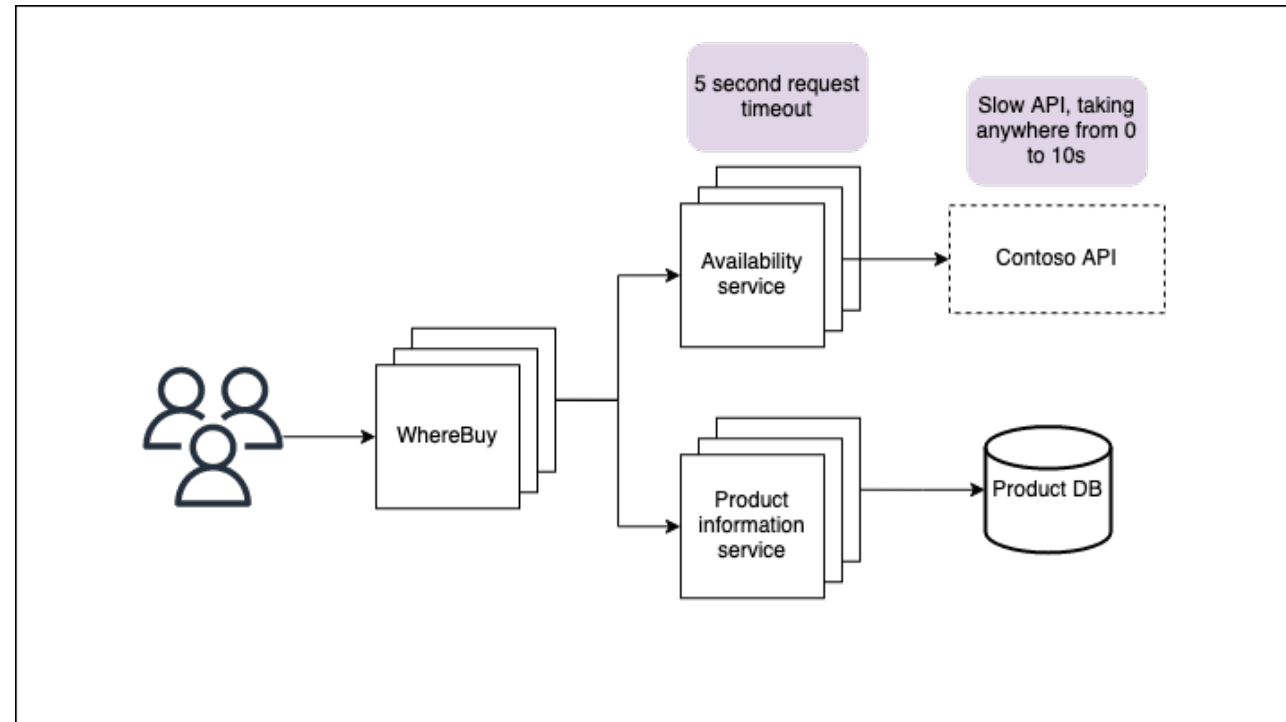
Chaos in the real world

Why microservices need to be resilient?

#wherebuy – low prices, more choices



Scenario A



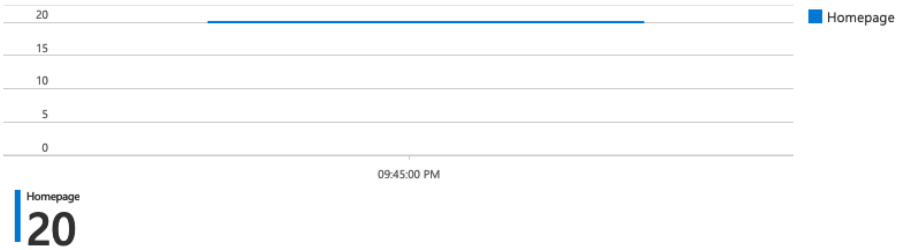
Scenario A

Load	Duration	Response time	Error percentage	Throughput
390	1 min, 7 secs	5.21 secs	40.00 %	5.91 /s
Total requests		90th percentile response time	Aggregate requests which failed	Request rate

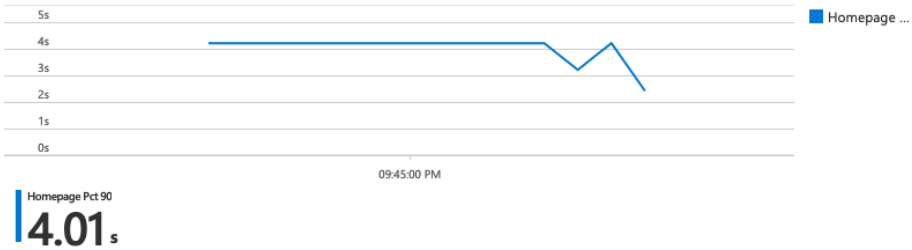
Client-side metrics

Requests : **Homepage** Percentile : **90** Error type : **2 selected** Time range : **10/6/2022, 9:44:08 PM - 10/6/2022, 9:45:16 PM** Group by : **5s**

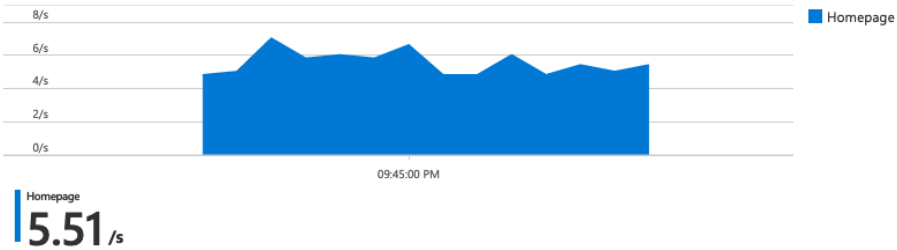
Virtual Users (Max)



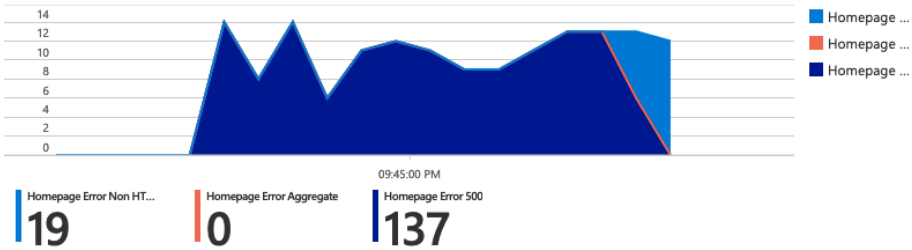
Response time (successful responses)



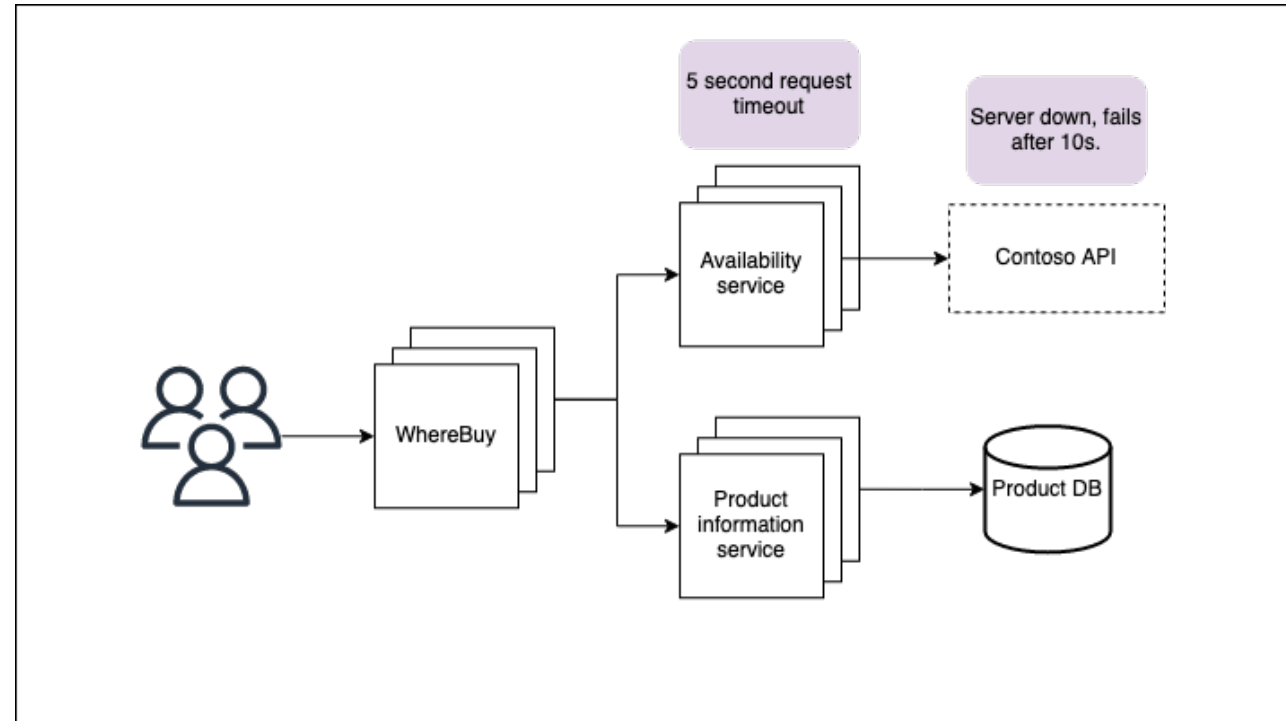
Requests/sec (Avg)



Errors (total)



Scenario B



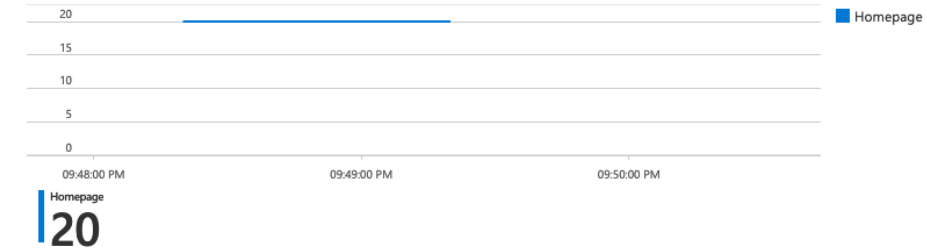
Scenario B

Load	Duration	Response time	Error percentage	Throughput
376	1 min, 1 sec	5.21 secs	100.00 %	6.16 /s
Total requests		90th percentile response time	Aggregate requests which failed	Request rate

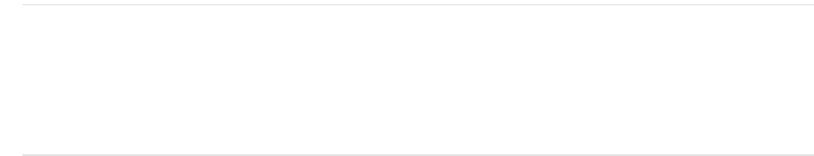
Client-side metrics

Requests : **Homepage** Percentile : **90** Error type : **2 selected** Time range : **10/6/2022, 9:47:58 PM - 10/6/2022, 9:49:11 PM** Group by : **5s**

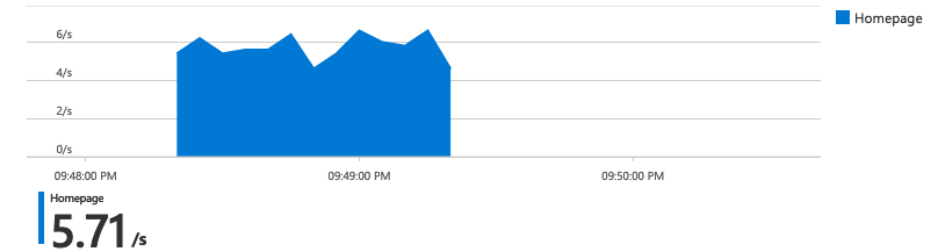
Virtual Users (Max)



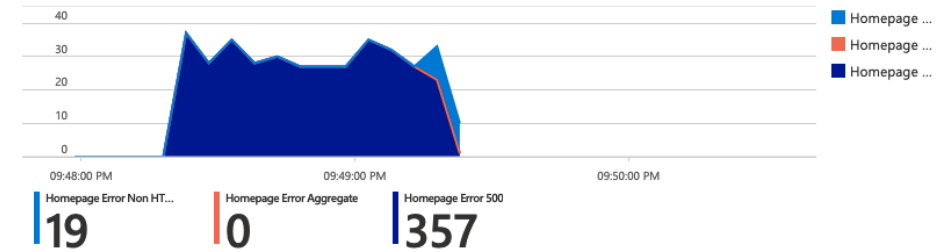
Response time (successful responses)



Requests/sec (Avg)



Errors (total)



Resiliency with Dapr

Defining and applying fault tolerance resiliency policies

Overview

How does a failure impact the overall service health and end users?

Overview

How does a failure impact the overall service health and end users?

Reminder: Failure is inevitable.

Overview

How does a failure impact the overall service health and end users?

Reminder: Failure is inevitable.

Dapr allows defining popular resiliency **policies**, which can be applied to **targets**, and **scoped** to specific applications.

Overview

How does a failure impact the overall service health and end users?


Reminder: Failure is inevitable.

Dapr allows defining popular resiliency **policies**, which can be applied to **targets**, and **scoped** to specific applications.

Declarative and decoupled from application code.

Policies - Timeouts

- Failures can be slow
- Upper-bound for waiting



```
spec:
  policies:
    timeouts:
      general: 5s
      largeResponse: 10s
      reallySlow: 60s
```

Policies - Retries

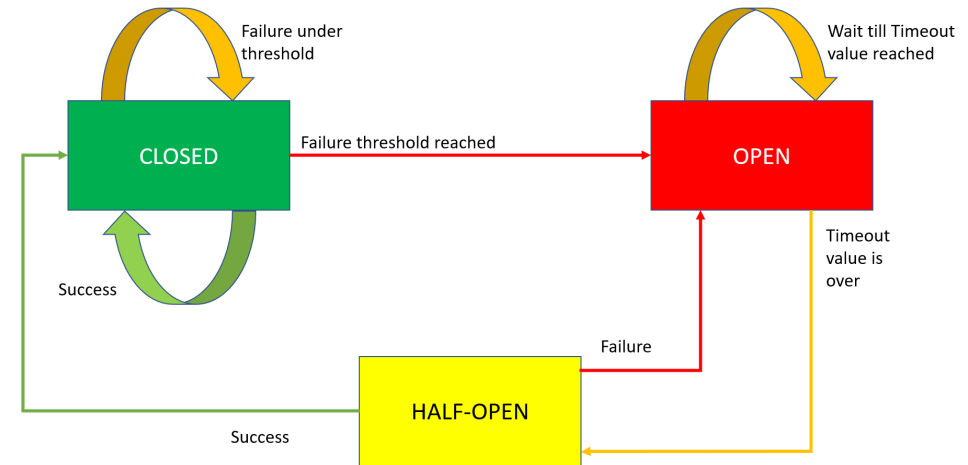
Choose between **constant** or **exponential** back-off retry policies



```
spec:
  policies:
    retries:
      retryFinite:
        policy: constant # Default policy
        duration: 5s # Time interval between retries
        maxRetries: 10
      retryForever:
        policy: exponential
        maxInterval: 15s
        maxRetries: -1 # Retry indefinitely (default)
```

Policies – Circuit breakers

Handle elevated failure rates by **shutting off all traffic** to the impacted service when a certain criteria is met.



Source: [\[Proposal\] Resiliency policies across all building blocks #3586 | dapr/dapr | GitHub](#)

```
spec:
  policies:
    circuitBreakers:
      ExternalApiCallerCB:
        maxRequests: 1 # maximum requests when CB is half-open
        interval: 8s # period for clearing internal counts
        timeout: 45s # period of open state before switching to half-open
        trip: consecutiveFailures > 8 # CEL statement to trigger open state
```


Targets

(1) Apps

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Target service invocation
calls between Dapr apps.

Targets

(1) Apps

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Target service invocation calls between Dapr apps.

(2) Components

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Target **inbound** (the sidecar calling your app) and **outbound** (calls to the sidecar) operations.

Targets

(1) Apps

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Target service invocation calls between Dapr apps.

(2) Components

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Target **inbound** (the sidecar calling your app) and **outbound** (calls to the sidecar) operations.

(3) Actors

- ✓ Retries
- ✓ Timeouts
- ✓ Circuit breakers

Circuit breaker can target an actor ID, all actors across the actor type, or both.

Dapr-izing your applications

Resilient microservices out-of-box.

Using Dapr with your application

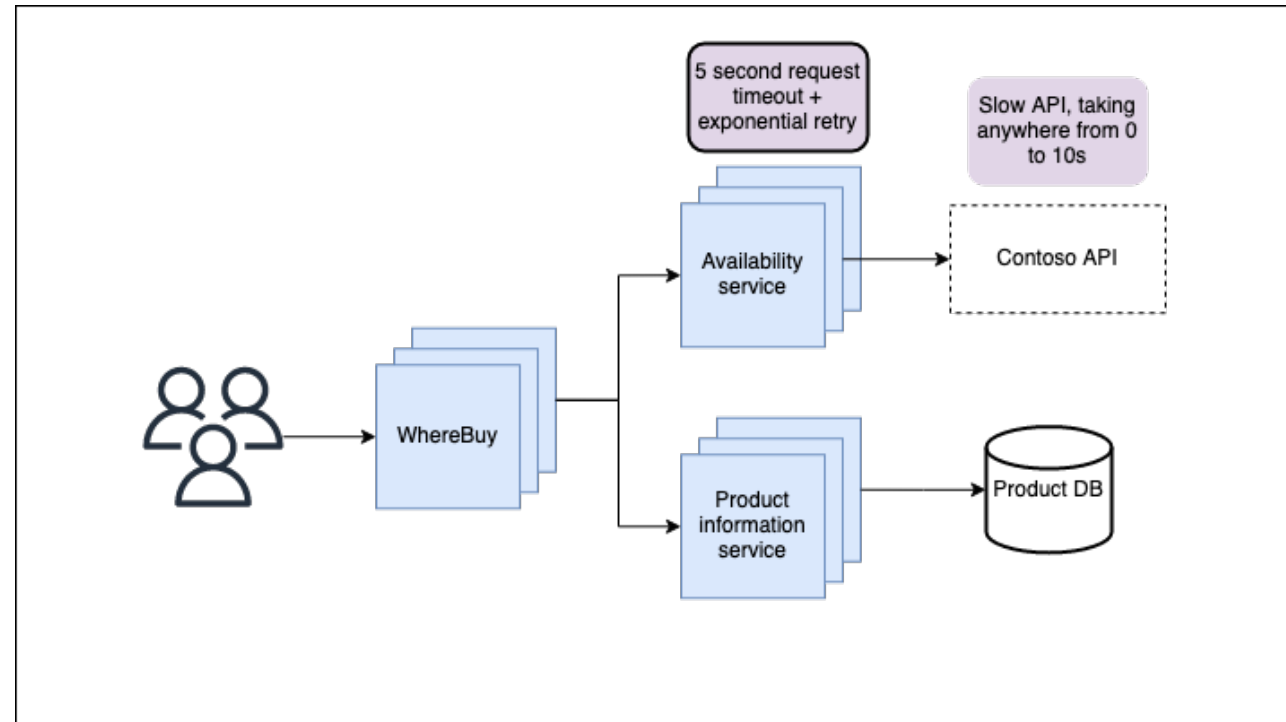
(1) Configure Dapr annotations

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: availabilityapp
  labels:
    app: availability
spec:
  template:
    metadata:
      # Just add these to enable Dapr!
      annotations:
        dapr.io/enabled: "true"
        dapr.io/app-id: "availabilityapp"
        dapr.io/enable-api-logging: "true"
        dapr.io/log-level: "debug"
    ...
```

(2) Invoke services using Dapr

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontendapp
  labels:
    app: frontend
...
spec:
  containers:
    - name: frontend
      image: ghcr.io/shubham1172/daprcon/frontend:latest
      env:
        # Use these when Dapr is disabled
        # - name: WHEREBUY_AVAILABILITY_API_URL
        #   value: http://availability/check
        # - name: WHEREBUY_PRODUCT_API_URL
        #   value: http://product/get
        # Use these when Dapr is enabled
        - name: WHEREBUY_AVAILABILITY_API_URL
          value: http://localhost:3500/v1.0/invoke/availabilityapp/method/check
        - name: WHEREBUY_PRODUCT_API_URL
          value: http://localhost:3500/v1.0/invoke/productapp/method/get
    ...
```

Scenario A – with Dapr



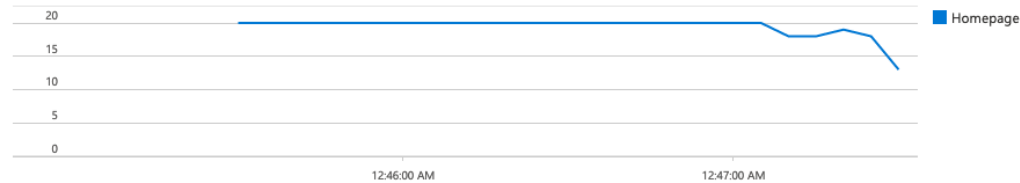
Scenario A – with Dapr

Load 448 Total requests	Duration 2 mins	Response time 11.40 secs 90th percentile response time	Error percentage 0.67 % Aggregate requests which failed	Throughput 3.76 /s Request rate
---	----------------------------------	--	---	---

Client-side metrics

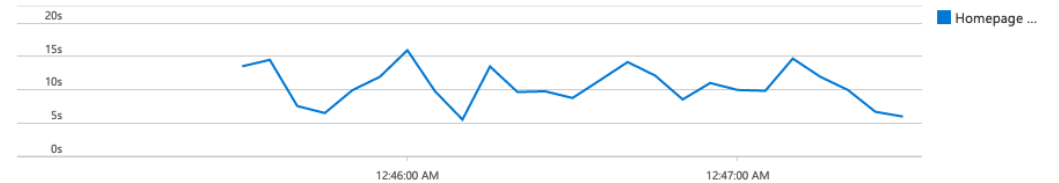
Requests : **Homepage** Percentile : **90** Error type : **500** Time range : **10/7/2022, 12:45:00 AM - 10/7/2022, 12:47:14 AM** Group by : **5s**

Virtual Users (Max)



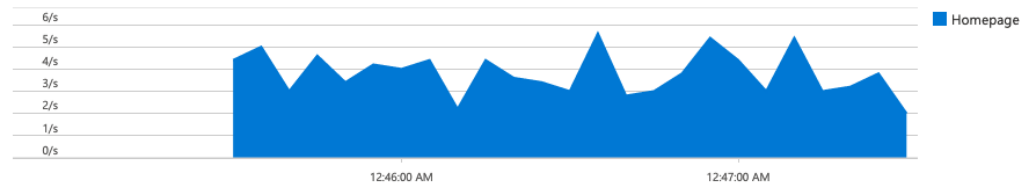
Homepage
20

Response time (successful responses)



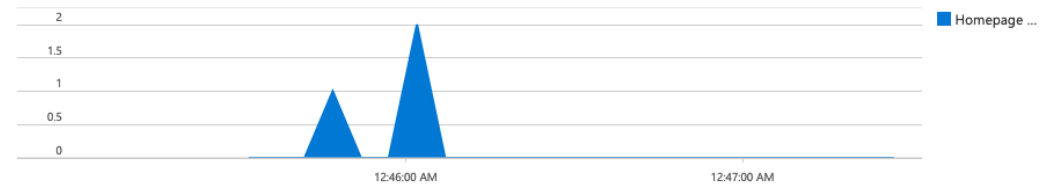
Homepage Pct 90
10.51s

Requests/sec (Avg)



Homepage
3.8 /s

Errors (total)



Homepage Error 500
3

Scenario A – comparison

Without Dapr

Requests/sec (Avg)



Homepage
5.51 /s

Error percentage

40.00 %

Aggregate requests which failed

With Dapr

Requests/sec (Avg)



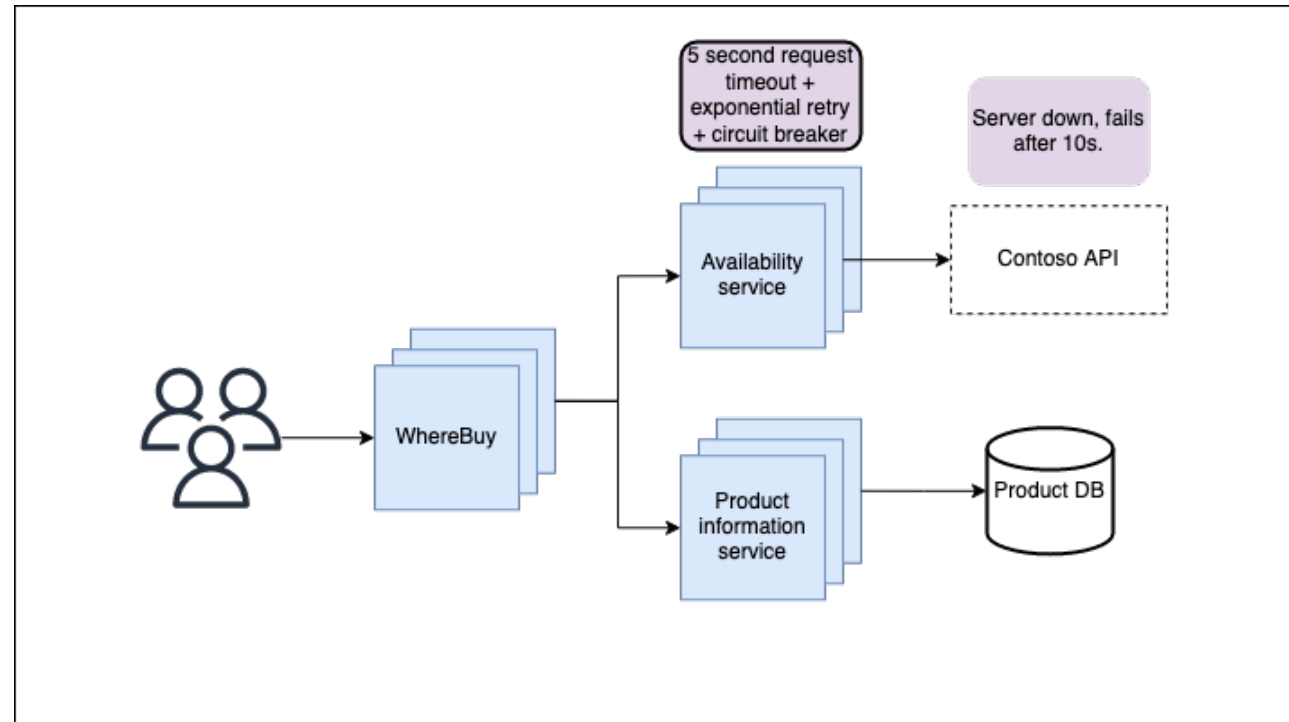
Homepage
3.8 /s

Error percentage

0.67 %

Aggregate requests which failed

Scenario B – with Dapr



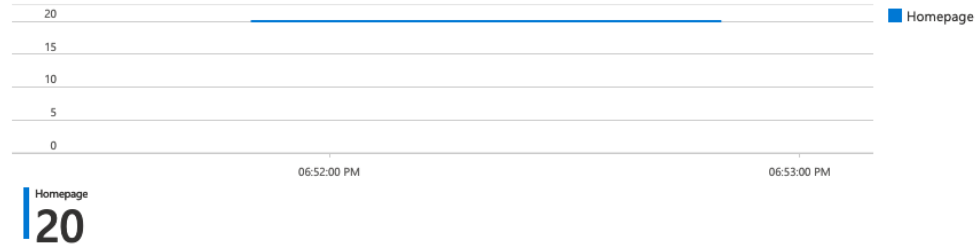
Scenario B – with Dapr

Load 5677 Total requests	Duration 1 min, 3 secs	Response time 205.00 ms 90th percentile response time	Error percentage 100.00 % Aggregate requests which failed	Throughput 91.56 /s Request rate
--	---	---	---	--

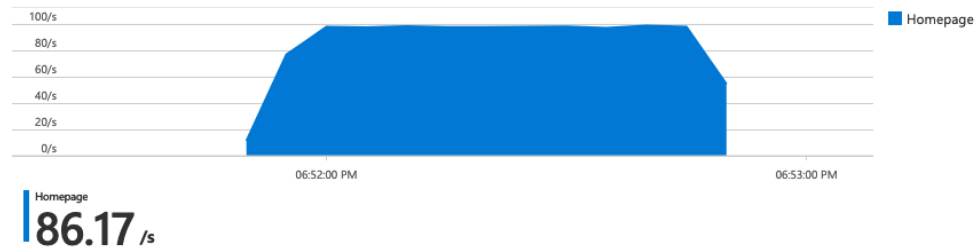
Client-side metrics

Requests : **Homepage** Percentile : **90** Error type : **2 selected** Time range : **10/7/2022, 6:51:27 PM - 10/7/2022, 6:52:34 PM** Group by : **5s**

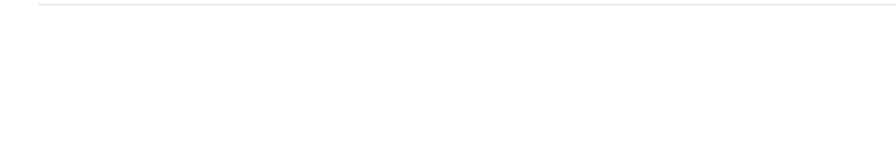
Virtual Users (Max)



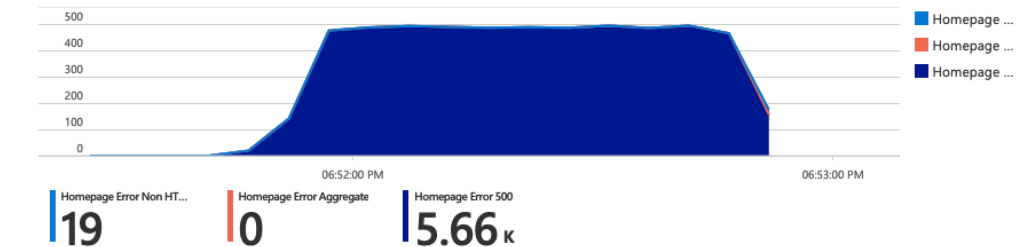
Requests/sec (Avg)



Response time (successful responses)

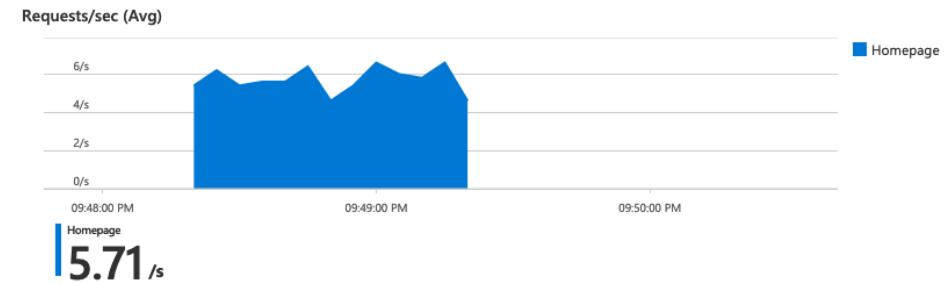
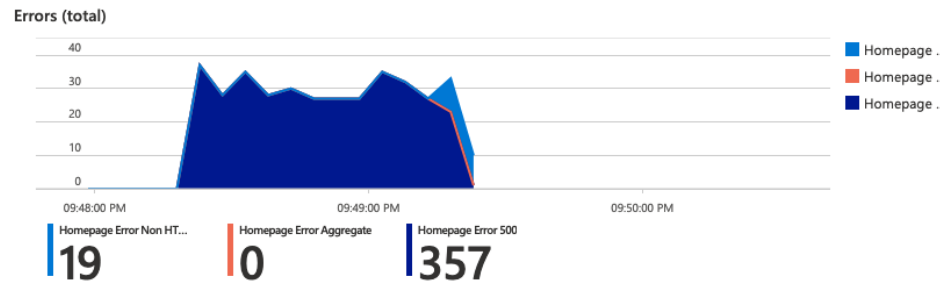


Errors (total)

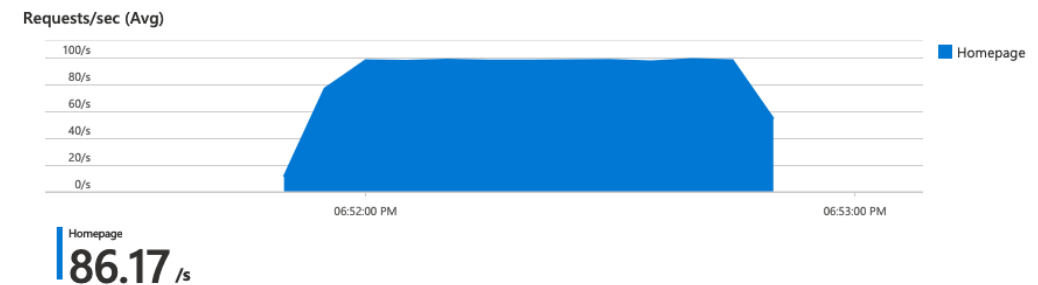
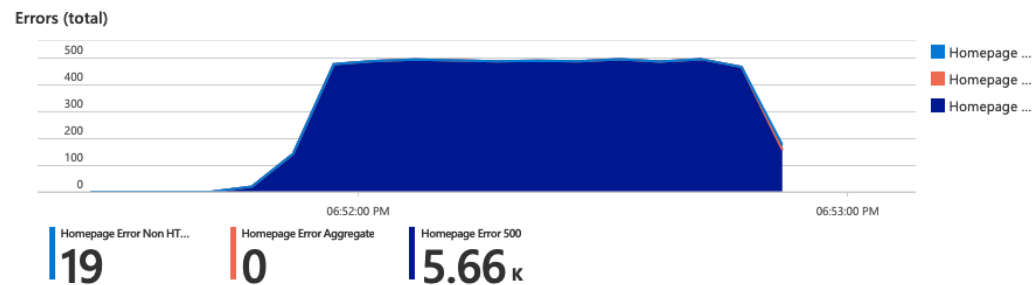


Scenario B – comparison

Without Dapr



With Dapr



Resources

Documentation

[Overview | Resiliency | Dapr Docs](#)

[Dapr arguments and annotations for daprd, CLI, and Kubernetes | Dapr Docs](#)

[\[Proposal\] Resiliency policies across all building blocks #3586](#)
[Original proposal on Dapr GitHub repository]

Code

[shubham1172/daprcon-resiliency](#)
[Code from this talk]



Thank you